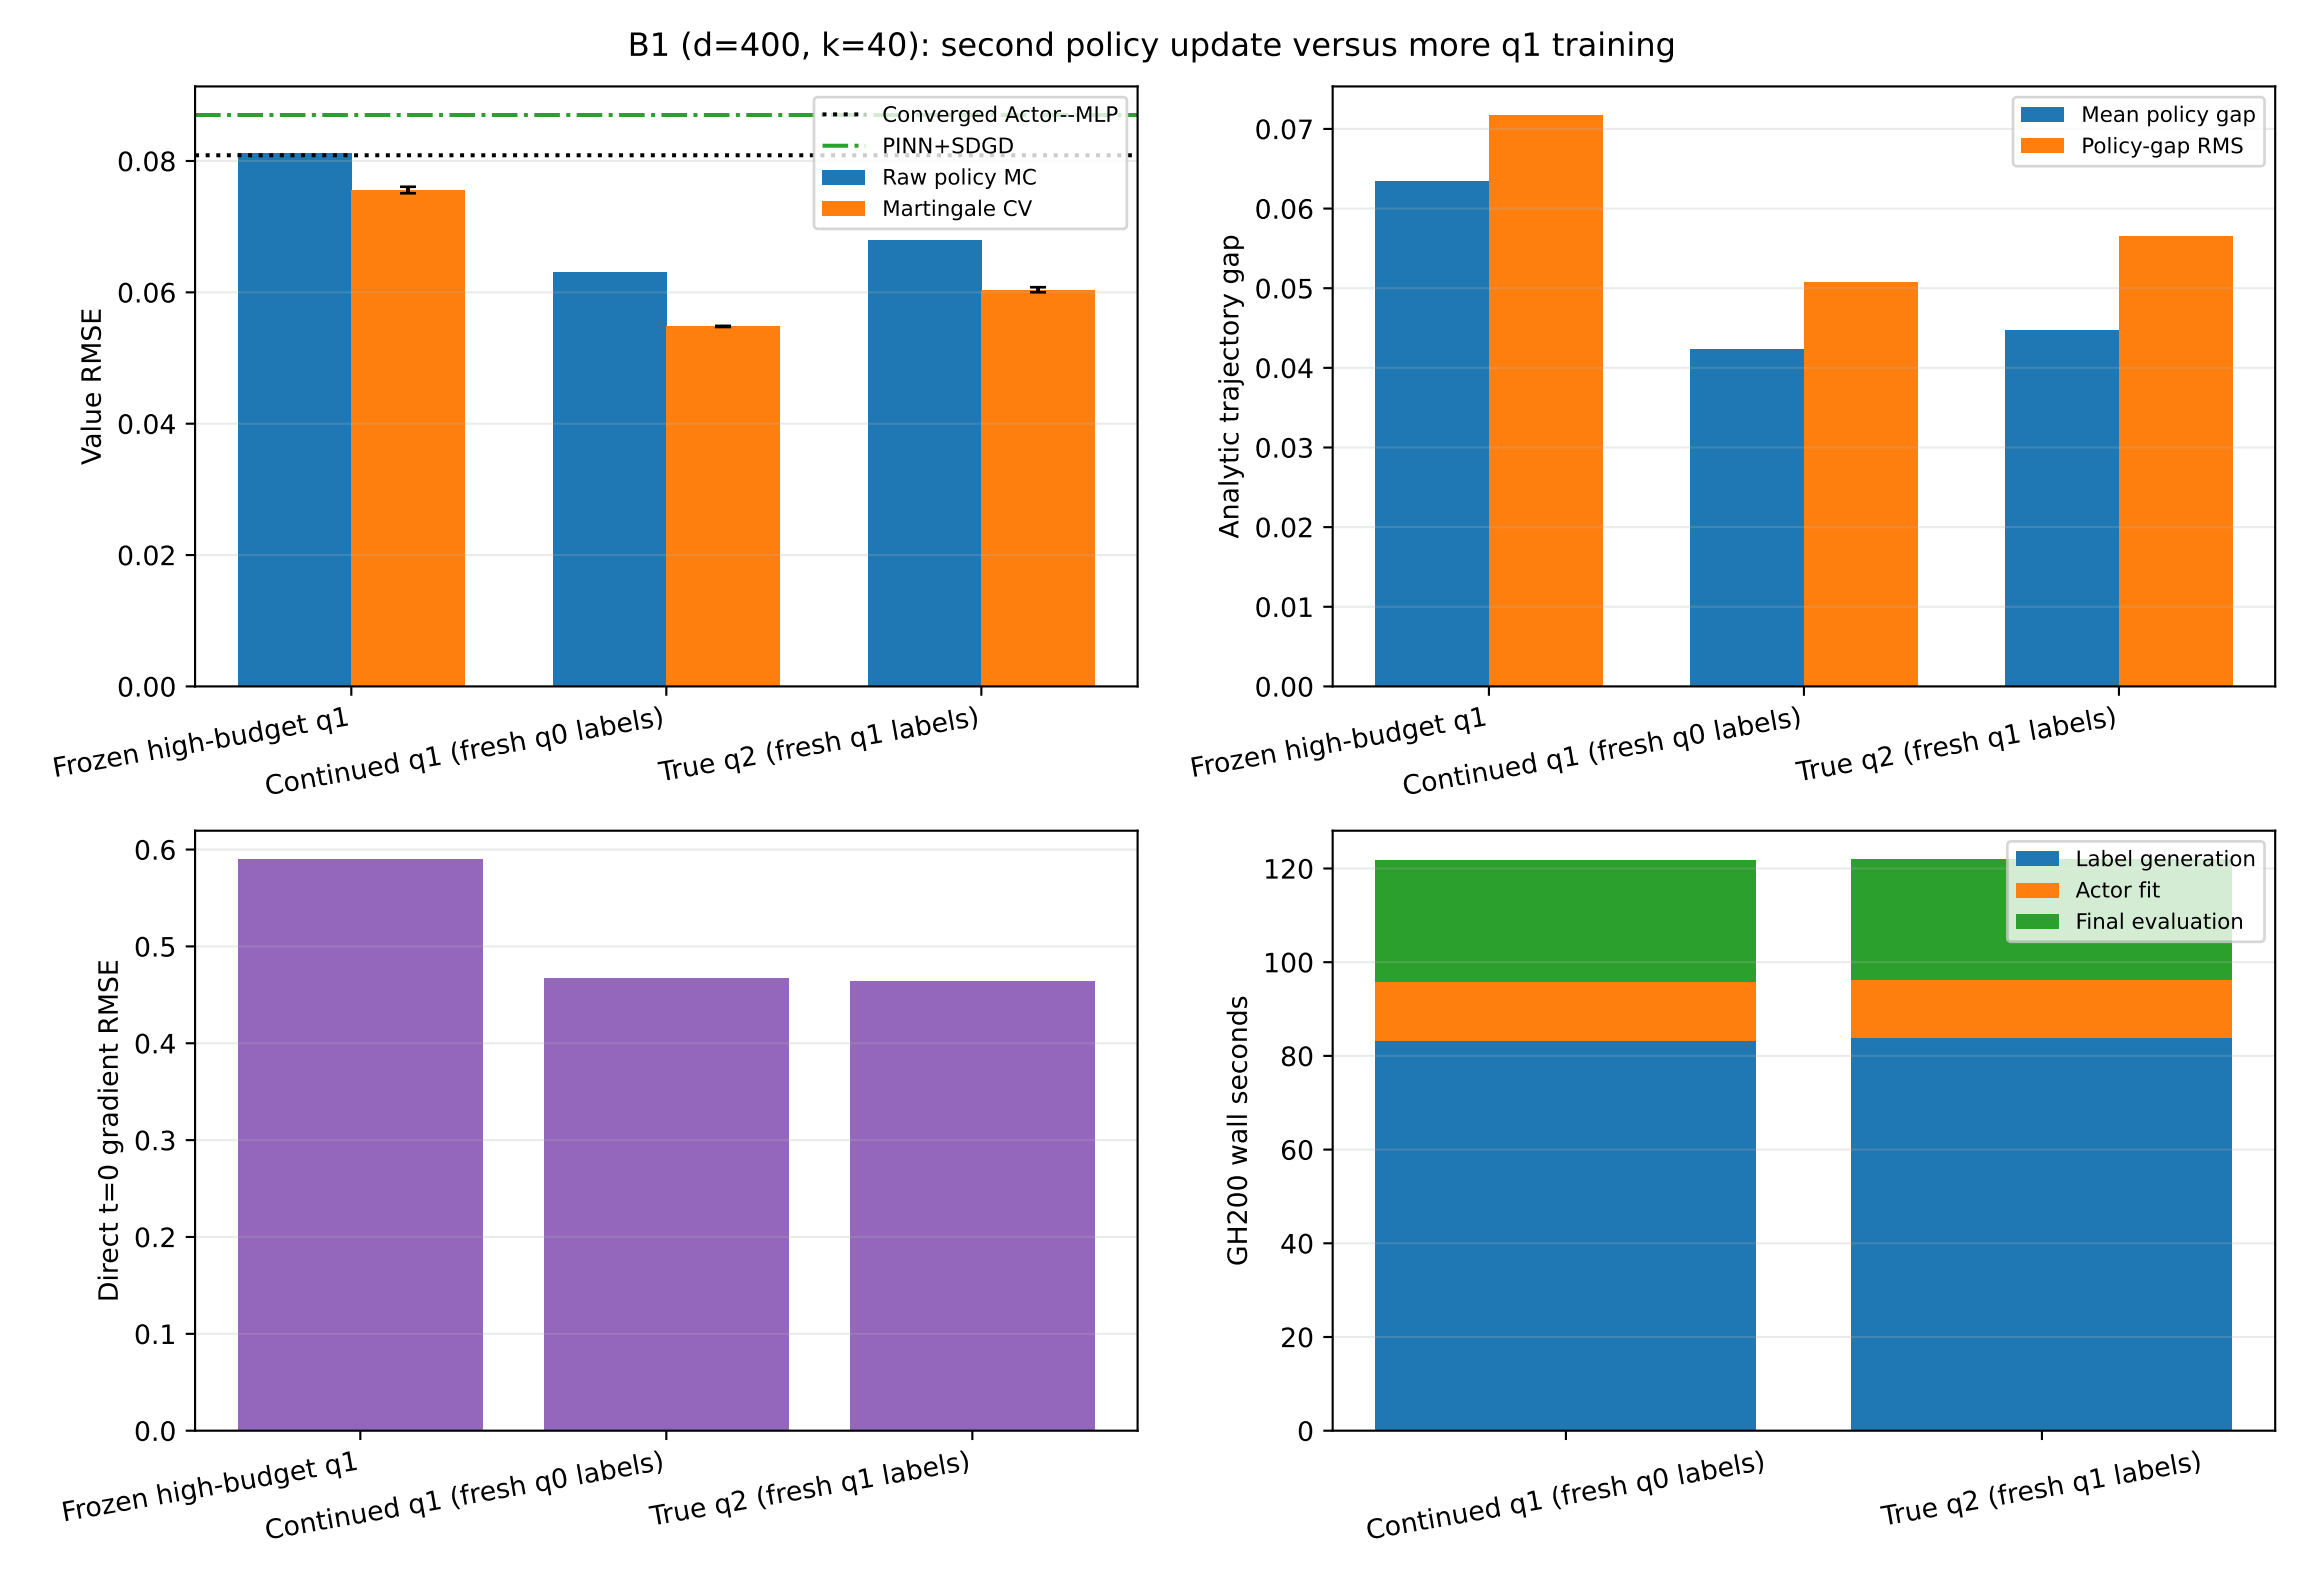
B1 (d=400, k=40): second policy update versus more q1 training
Converged Actor--MLP
PINN+SDGD
Raw policy MC
Martingale CV
Value RMSE
0.00
0.02
0.04
0.06
0.08
Frozen high-budget q1
Continued q1 (fresh q0 labels)
True q2 (fresh q1 labels)
Mean policy gap
Policy-gap RMS
Analytic trajectory gap
0.00
0.01
0.02
0.03
0.04
0.05
0.06
0.07
Frozen high-budget q1
Continued q1 (fresh q0 labels)
True q2 (fresh q1 labels)
Direct t=0 gradient RMSE
0.0
0.1
0.2
0.3
0.4
0.5
0.6
Frozen high-budget q1
Continued q1 (fresh q0 labels)
True q2 (fresh q1 labels)
Label generation
Actor fit
Final evaluation
GH200 wall seconds
0
20
40
60
80
100
120
Continued q1 (fresh q0 labels)
True q2 (fresh q1 labels)