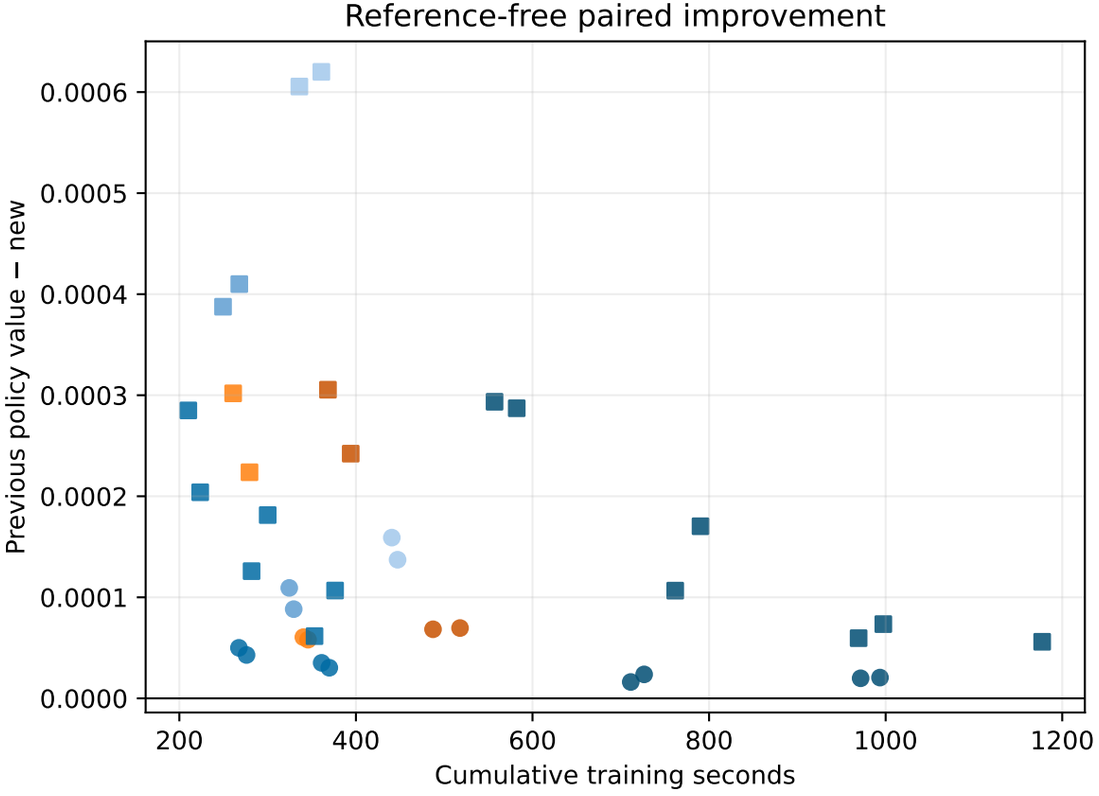
Reference-free paired improvement
Previous policy value − new
Cumulative training seconds
0.0006
0.0005
0.0004
0.0003
0.0002
0.0001
0.0000
200
400
600
800
1000
1200

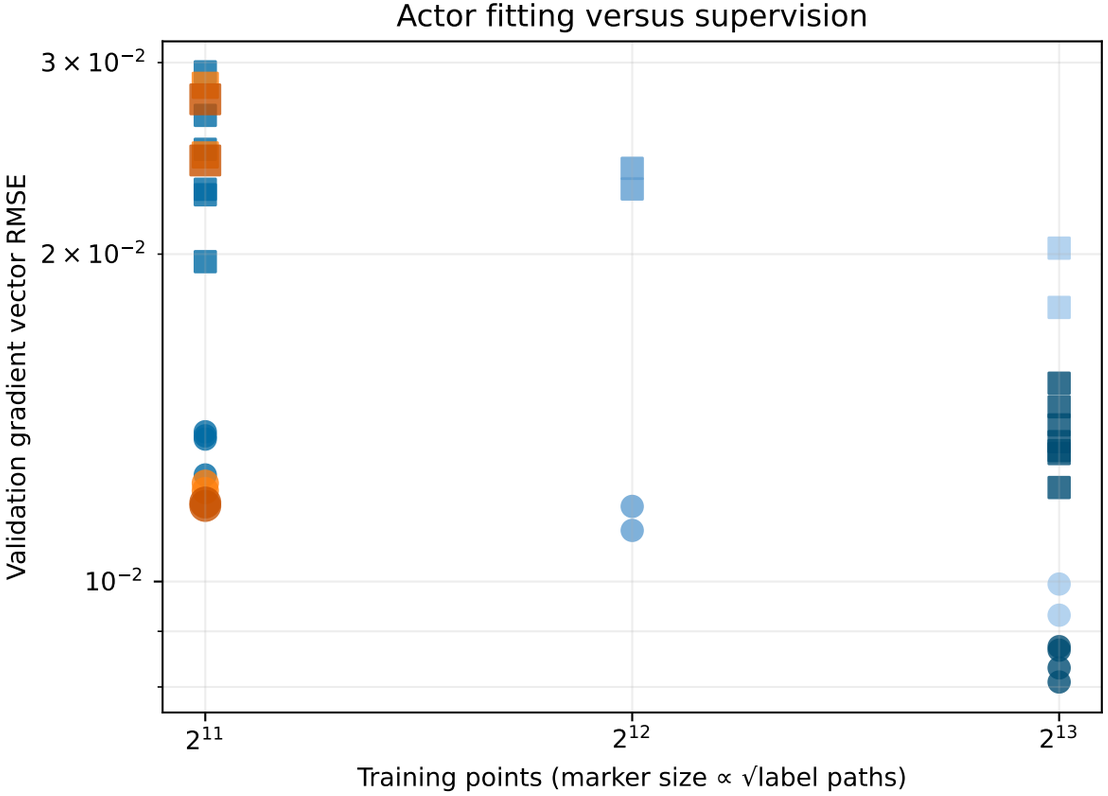
Actor fitting versus supervision
Validation gradient vector RMSE
Training points (marker size ∝ √label paths)
3 × 10⁻²
2 × 10⁻²
10⁻²
2¹¹
2¹²
2¹³

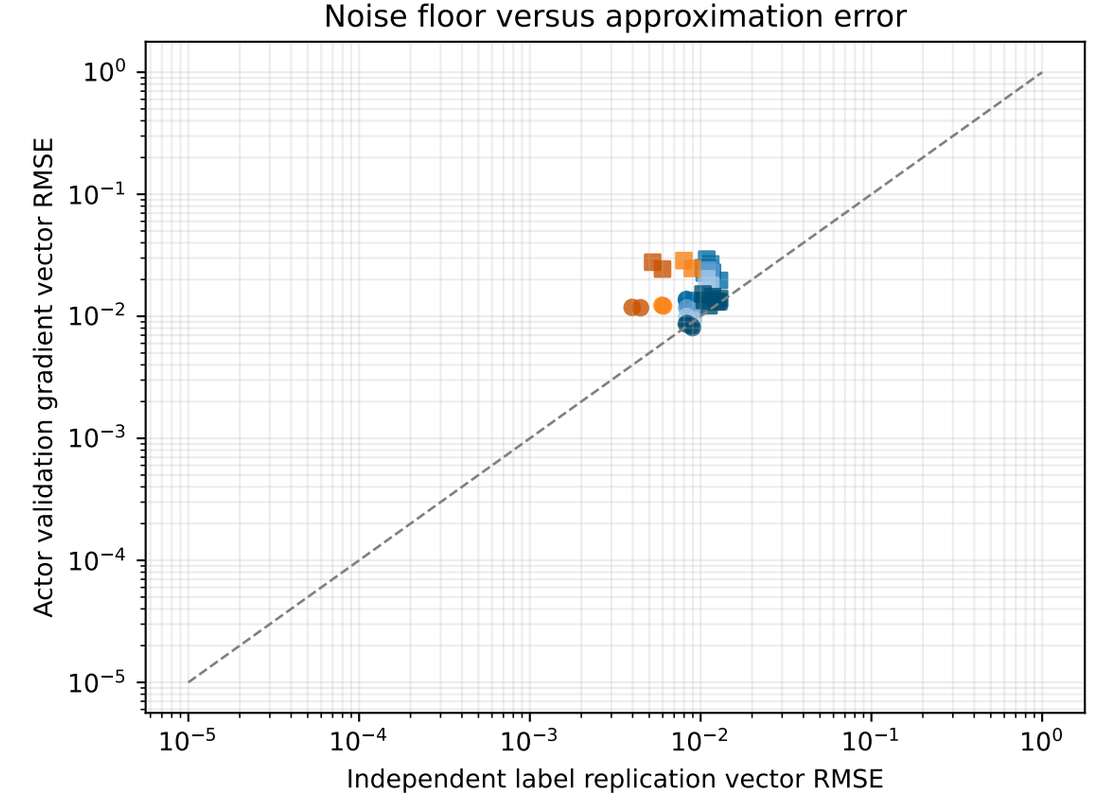
Noise floor versus approximation error
Actor validation gradient vector RMSE
Independent label replication vector RMSE
10⁰
10⁻¹
10⁻²
10⁻³
10⁻⁴
10⁻⁵

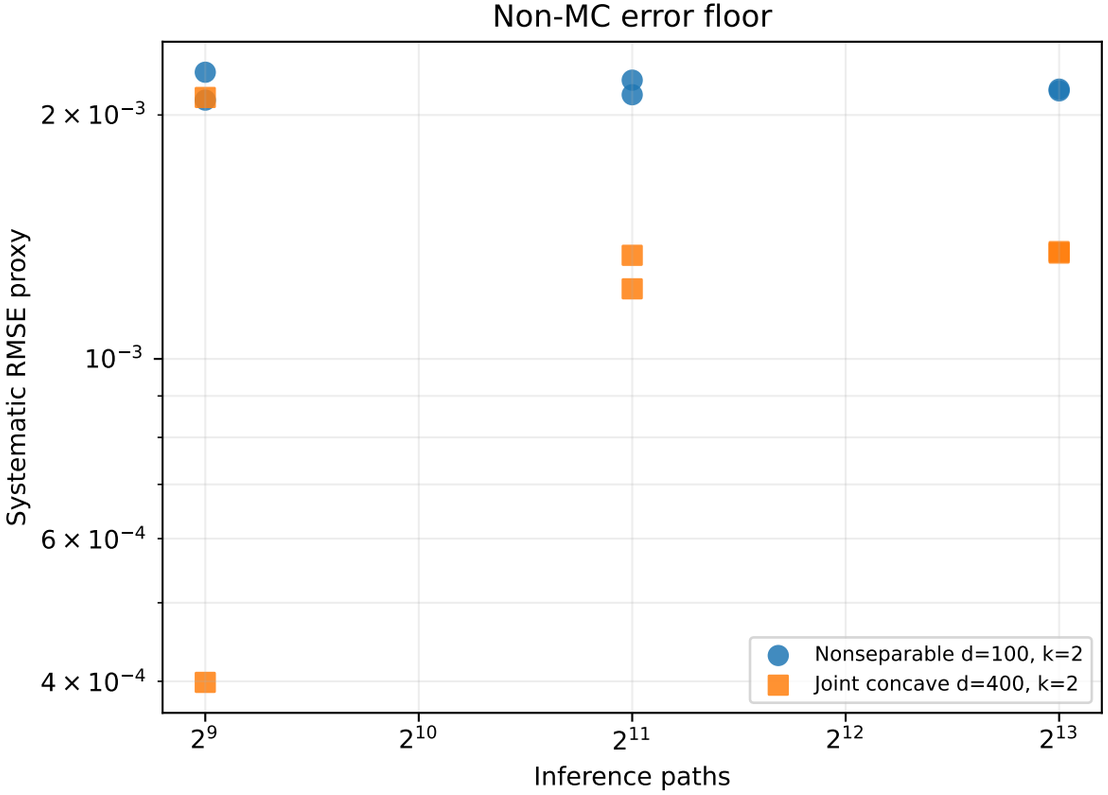
Non-MC error floor
Systematic RMSE proxy
Inference paths
2 × 10⁻³
10⁻³
6 × 10⁻⁴
4 × 10⁻⁴
2⁹
2¹⁰
2¹¹
2¹²
2¹³
Nonseparable d=100, k=2
Joint concave d=400, k=2